

## 2. Introduce scikit-learn as a machine learning library.

```
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn import metrics
iris = load_iris()
X = iris.data
y = iris.target
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
model = KNeighborsClassifier(n_neighbors=3)
model.fit(X_train, y_train)
predictions = model.predict(X_test)
accuracy = metrics.accuracy_score(y_test, predictions)
print(f'Accuracy: {accuracy}')
```

Output :-

Accuracy: 0.9777777777777777

## 3. Install and set up scikit-learn and other necessary tools. Procedure.

```
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
iris = datasets.load_iris()
X_train, X_test, y_train, y_test = train_test_split(iris.data, iris.target,
test_size=0.3)
clf = KNeighborsClassifier(n_neighbors=3)
clf.fit(X_train, y_train)
accuracy = clf.score(X_test, y_test)
print(f'Accuracy: {accuracy}')
```

Output :-

Accuracy: 0.9777777777777777

## 4. Write a program to Load and explore the dataset of .CSV and excel files using pandas.

```
import pandas as pd
csv_data = pd.read_csv('california_housing_test.csv')
print("First few rows of CSV file:")
print(csv_data.head())
print("\nSummary statistics of CSV file:")
print(csv_data.describe())
print("\nInformation about columns in CSV file:")
print(csv_data.info())
```

```
excel_data = pd.read_excel('SampleData.xlsx')
```

```

print("First few rows of Excel file:")
print(excel_data.head())
print("\nSummary statistics of Excel file:")
print(excel_data.describe())
print("\nInformation about columns in Excel file:")
print(excel_data.info())

```

Output :-

```

First few rows of CSV file:
longitude latitude housing_median_age total_rooms total_bedrooms \
0 -122.85 37.37 27.0 3885.0 661.0
1 -118.30 34.26 43.0 1510.0 310.0
2 -117.81 33.78 27.0 3589.0 507.0
3 -118.35 33.82 28.0 67.0 15.0
4 -119.67 36.33 19.0 1241.0 244.0

population households median_income median_house_value
0 1537.0 886.0 6.6885 344700.0
1 809.0 277.0 3.5990 176500.0
2 1484.0 495.0 5.7934 270500.0
3 49.0 11.0 6.1359 330000.0
4 850.0 237.0 2.9375 81700.0

Summary statistics of CSV file:
count longitude latitude housing_median_age total_rooms \
mean -119.582000 35.63839 28.845333 2599.578667
std 1.904936 2.1267 12.555396 2155.553312
min -124.180000 32.56000 1.000000 6.000000
25% -121.810000 33.93000 18.000000 1401.000000
50% -118.485000 34.27000 29.000000 2105.000000
75% -115.020000 37.69000 37.000000 3129.000000
max -114.490000 41.92000 52.000000 38450.000000

total_bedrooms population households median_income \
count 3000.000000 3000.000000 3000.000000 3000.000000
mean 529.950667 1402.798667 489.01200 3.807272
std 415.654368 1030.543012 365.42271 1.854512
min 2.000000 5.000000 2.00000 0.499900
25% 291.000000 700.000000 273.00000 2.544000
50% 437.000000 1155.000000 409.50000 3.487150
75% 636.000000 1742.750000 597.25000 4.656475
max 5419.000000 11935.000000 4930.00000 15.000100

median_house_value
count 3000.00000
mean 205846.27500
std 113119.68747
min 22500.00000
25% 121200.00000
50% 177650.00000
75% 263975.00000
max 500001.00000

Information about columns in CSV file:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3000 entries, 0 to 2999
Data columns (total 9 columns):
# Column Non-Null Count Dtype
---
0 longitude 3000 non-null float64
1 latitude 3000 non-null float64
2 housing_median_age 3000 non-null float64
3 total_rooms 3000 non-null float64
4 total_bedrooms 3000 non-null float64
5 population 3000 non-null float64
6 households 3000 non-null float64
7 median_income 3000 non-null float64
8 median house value 3000 non-null float64
dtypes: float64(9)
memory usage: 211.1 KB
None

First few rows of Excel file:
Unnamed: 0 Unnamed: 1 Unnamed: 2
0 NaN NaN NaN
1 NaN NaN NaN
2 NaN NaN NaN
3 NaN NaN Online Instruction Page
4 NaN NaN Sample Data for Excel

Summary statistics of Excel file:
count Unnamed: 0
mean NaN
std NaN
min NaN
25% NaN
50% NaN
75% NaN
max NaN

Information about columns in Excel file:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 26 entries, 0 to 25
Data columns (total 3 columns):
# Column Non-Null Count Dtype
---
0 Unnamed: 0 0 non-null float64
1 Unnamed: 1 3 non-null object
2 Unnamed: 2 20 non-null object
dtypes: float64(1), object(2)
memory usage: 756.0+ bytes
None

```

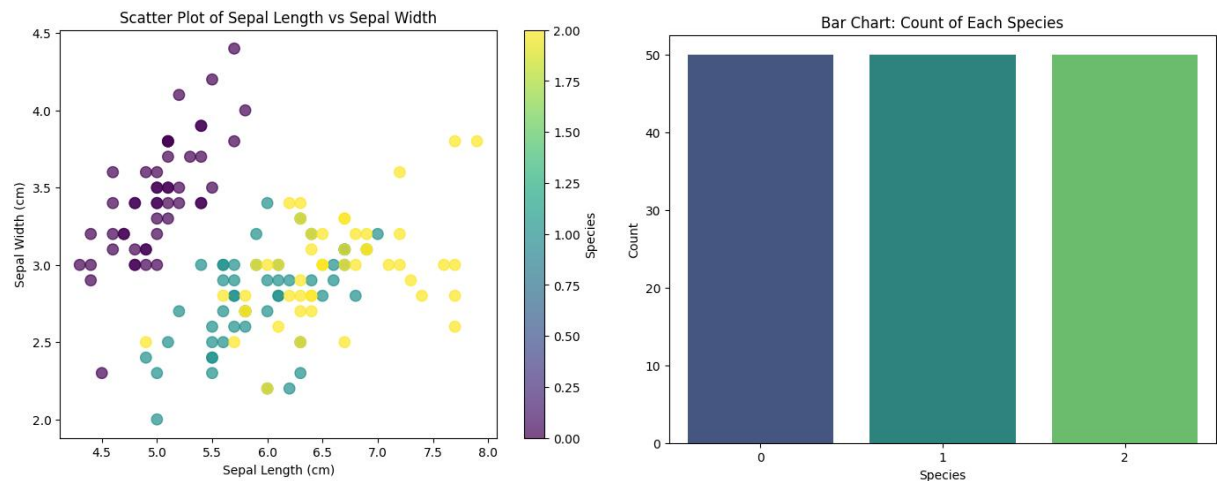
5. Write a program to Visualize the dataset to gain insights using Matplotlib or Seaborn by plotting scatter plots, bar charts.

```

import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.datasets import load_iris
import pandas as pd
iris = load_iris()
iris_df = pd.DataFrame(data=iris.data, columns=iris.feature_names)
iris_df['target'] = iris.target
plt.figure(figsize=(8, 6))
plt.scatter(iris_df['sepal length (cm)'], iris_df['sepal width (cm)'],
c=iris_df['target'], cmap='viridis', s=80, alpha=0.7)
plt.xlabel('Sepal Length (cm)')
plt.ylabel('Sepal Width (cm)')
plt.title('Scatter Plot of Sepal Length vs Sepal Width')
plt.colorbar(label='Species')
plt.show()
plt.figure(figsize=(8, 6))
sns.countplot(x='target', data=iris_df, palette='viridis')
plt.xlabel('Species')
plt.ylabel('Count')
plt.title('Bar Chart: Count of Each Species')
plt.show()

```

Output :-



6. Write a program to implement a k-Nearest Neighbours (k-NN) classifier using scikitlearn and Train the classifier on the dataset and evaluate its performance.

```
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn import metrics
iris = load_iris()
X = iris.data
y = iris.target
# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
random_state=42)
k = 3
knn_classifier = KNeighborsClassifier(n_neighbors=k)
knn_classifier.fit(X_train, y_train)
predictions = knn_classifier.predict(X_test)
accuracy = metrics.accuracy_score(y_test, predictions)
print(f'Accuracy: {accuracy}')
print("Classification Report:")
print(metrics.classification_report(y_test, predictions))
print("Confusion Matrix:")
print(metrics.confusion_matrix(y_test, predictions))
```

Output :-

```
Accuracy: 1.0
Classification Report:
              precision    recall  f1-score   support

     0         1.00      1.00      1.00        19
     1         1.00      1.00      1.00        13
     2         1.00      1.00      1.00        13

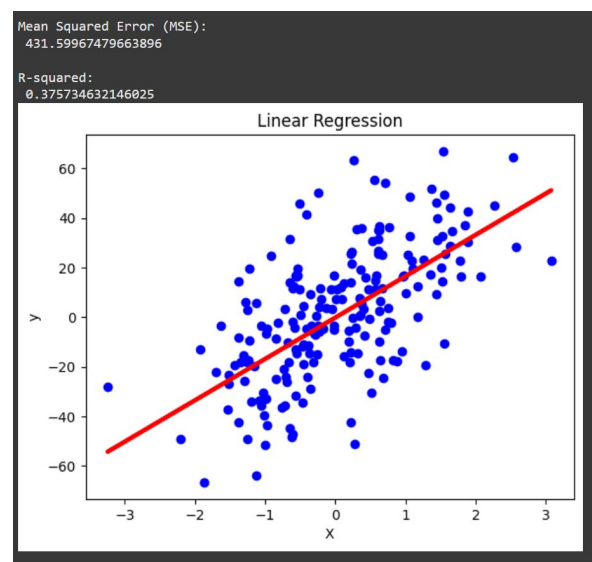
   accuracy                   1.00        45
  macro avg              1.00      1.00      1.00        45
 weighted avg              1.00      1.00      1.00        45

Confusion Matrix:
[[19  0  0]
 [ 0 13  0]
 [ 0  0 13]]
```

7. Write a program to implement a linear regression model for regression tasks and Train the model on a dataset with continuous target variables.

```
import numpy as np
from sklearn.linear_model import LinearRegression
from sklearn.datasets import make_regression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score
import matplotlib.pyplot as plt
X,y = make_regression(n_samples=1000, n_features=1, noise=20, random_state=42)
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.2,
random_state=42)
linear_reg = LinearRegression()
linear_reg.fit(X_train, y_train)
predictions = linear_reg.predict(X_test)
mse = mean_squared_error(y_test, predictions)
r2 = r2_score(y_test, predictions)
print("Mean Squared Error (MSE):\n",mse)
print("R-squared:\n",r2)
plt.scatter(X_test, y_test, color='blue')
plt.plot(X_test, predictions, color='red', linewidth=3)
plt.xlabel('X')
plt.ylabel('y')
plt.title('Linear Regression')
plt.show()
```

Output :-



8. Write a program to implement a decision tree classifier using scikitlearn and visualize the decision tree and understand its splits.

```
from sklearn.datasets import load_iris
from sklearn.tree import DecisionTreeClassifier, plot_tree
import matplotlib.pyplot as plt
```

```

iris = load_iris()
X = iris.data
y = iris.target
class_names = [str(name) for name in iris.target_names]
decision_tree = DecisionTreeClassifier()
decision_tree.fit(X, y)
plt.figure(figsize=(12, 8))
plot_tree(decision_tree, feature_names=iris.feature_names,
class_names=class_names, filled=True, rounded=True)
plt.title("Decision Tree Visualization")
plt.show()

```

Output :-

